

THE PRIMER • PART 2

Where your draft actually goes.

You paste three paragraphs into a chat box to tidy the language. This is what happens to them — and the one setting that changes the answer.

This issue assumes you know nothing about how these systems work. Every term is explained on the page, before it is used. Part 1 explained what the thing inside is. This part is about where your words go once you hand them over.

Rushikesh R. Mahajan, Advocate

Data protection, technology and AI governance • LL.M. in Law and Technology, Queen's University Belfast
Practises before the Bombay High Court and courts across India in civil and criminal matters

What this is / is not. This is general, illustrative and **non-exhaustive** information, written for lawyers, about what happens to text typed into a commercial AI product. It is **not legal advice**, **not a compliance determination**, and does **not** create an advocate-client relationship. Vendor terms are quoted as published on the dates given and change without notice — re-read the page before you rely on it. Readers outside India should map every term to their own law. Current as at 13 September 2026.

Prepared by Rushikesh Mahajan, Advocate.

Start here

It is a Tuesday. You are drafting a reply. One paragraph is not coming out cleanly, so you open ChatGPT or Claude in a browser tab, paste three paragraphs of your own draft in, and ask it to tighten the language.

Those three paragraphs contain your client's name. The disputed sum. The argument you intend to run.

Where did they just go?

That is the whole subject of this issue, and it is worth being clear at the outset about why it is *this* question and not the more famous one.

Almost everything written about AI and privacy is about **training** — whether these companies were entitled to use the world's books, articles and websites to build their products in the first place. That is a real and serious question. **It is also not your problem.** You were not consulted when those models were built and you have no exposure from it.

Your problem is the three paragraphs you pasted in on Tuesday. That is a different event, governed differently, and almost nobody writes about it.

Why Part 1 had to come first

Part 1 of this series made a single claim: **a language model predicts the next word, and that is the whole machine.** Given some text, it produces a guess at the next fragment, picks one, adds it on, and repeats. There is no step in that loop where it consults a source, checks a fact, or decides what is true.

In Part 1 that looked like a point about *accuracy* — why these systems invent citations, and why the verification burden never moves. It is. But it carries a second consequence, and that consequence is the foundation of everything in this issue.

THE CONSEQUENCE PART 1 SETS UP

If the model is only a next-word predictor, then it holds nothing. No filing cabinet. No record of your conversation. No memory of you between Tuesday and Wednesday. It reads what it is handed, produces an answer, and retains nothing at all.

Which forces the question this issue exists to answer.

Your conversation history is obviously kept somewhere — you can scroll up and read it tomorrow. The thing you are using does appear to remember your name, your preferences, the document you uploaded last week. **If none of that is inside the model, then all of it is somewhere else** — somewhere with a different owner, a different physical location, and a different set of legal duties attached to it.

Part 1 emptied the box. Part 2 is about where the contents actually are.

There is a second thing Part 1 gives us. It established that these models are **built** by being shown enormous quantities of text, and that the building is what produced the behaviour. That is what makes “*was my text used to build it?*” a sensible question at all — and, as we will see, a question with four different answers depending on which account you happened to be logged into.

So the two parts do different jobs, and they only work in this order. **Part 1 told you what the thing is. Part 2 tells you where your words go.** You cannot answer the second without the first, because almost every wrong answer in this area starts by imagining that “the AI” is one thing that holds your data — when the one part properly called the model is the one part that holds nothing.

Part 1 — *It predicts the next word. That is the whole machine.* Free in the archive: buttondown.com/the_desk/archive/

Two nouns, and two events

Most confusion here comes from one habit: people use “AI”, “the model”, “the machine” and “the system” as though they were the same thing. They are not. This issue uses exactly **two nouns** — and keeps apart **two events** that are constantly confused. Here they are, before anything is built on them.

THE MODEL

The thing Part 1 described. Given some text, it produces a guess at what word comes next, then does that again, and again. That is all it does.

It has no memory. It cannot look anything up. It cannot send an email or open a file. On its own it is inert — closer to a very elaborate autocomplete than to an assistant.

THE PRODUCT

The thing you actually log into and pay for: ChatGPT, Claude, Copilot, Gemini. **The model is one component inside it.** The product is the model *plus* the screen, plus the memory, plus the search and file tools, plus the filters — built and run by a company, often with the help of other companies.

When someone says “we use ChatGPT”, they are describing **the product**. Nearly every question in this issue turns on the difference between the two.

And two events, which are constantly mistaken for each other:

TRAINING — HOW THE PRODUCT WAS BUILT

Text is used to **build** the model, before you ever see it. This happened in the past, at the vendor, at enormous scale.

USE — WHAT HAPPENS WHEN YOU TYPE

You send text **to** a product that already exists. It travels through that product, is worked on, and an answer comes back. This happens every time you press enter.

Keep those apart and most of the fog clears. **A vendor who tells you “we do not train on your data” has made a statement about the first event and said nothing at all about the second.** Both things can be true at the same time: your paragraphs were never used to build anything, *and* your paragraphs passed through several companies in several countries on Tuesday afternoon.

Why training is hard to undo — plainly

It is worth one short detour, because the reason deletion is difficult here is genuinely misunderstood, and it is not obstruction.

A model is, physically, a very long list of numbers. Billions of them. Those numbers are usually called *parameters*, or *weights*. Training is the process of nudging those numbers up and down, over and over, until the model’s guesses about the next word get good.

Nothing is stored as text. There is no folder inside the model containing your sentences. There is only the effect your sentences had on the numbers — mixed in with the effect of everything else it ever saw.

WHAT THAT MEANS FOR A DELETION REQUEST

In an ordinary database you can delete a record. There is a row, you find it, it goes.

Here there is no row. Your sentence is not sitting in the model as a sentence, so there is nothing to locate and nothing to strike out. You cannot subtract one book from a person’s education. Removing the influence generally means rebuilding the model — which is why vendors resist it, and why “just delete it” is a harder instruction than it sounds.

This is also why **the tier you are on matters so much**: if your text never enters training in the first place, this problem never arises for you.

European regulators have already looked at this. In **Opinion 28/2024**, adopted on 17 December 2024, the European Data Protection Board held that a model trained on personal data cannot simply be *assumed* to be anonymous — it has to be assessed case by case, because personal data may remain embedded in the model’s parameters.

Source: EDPB Opinion 28/2024 — edpb.europa.eu/our-work-tools/our-documents/opinion-board-art-64/opinion-282024-certain-data-protection-aspects_en · full text (PDF): edpb.europa.eu/system/files/2024-12/edpb_opinion_202428_ai-models_en.pdf

What you type, and who keeps it

Now the practical question. You pasted three paragraphs in. Will they be used to build the next version of the product?

The honest answer is: it depends entirely on which account you were logged into. Same company, same product name, same screen — opposite answer. This is the single most useful thing in this issue, and it is checkable in two minutes on the vendors’ own pages.

What you are logged into	Used to train the model?	The vendor's own words
ChatGPT — Free, Plus, Pro (individual accounts)	YES, unless you opt out	"When you use our services for individuals such as ChatGPT and Codex, we may use your content to train our models. " Opt out at Settings → Data Controls → "Improve the model for everyone".
ChatGPT — Business, Enterprise, Edu; and the API	No, by default	"By default, we do not train on any inputs or outputs from our products for business users , including ChatGPT Business, ChatGPT Enterprise, and the API."
Claude — Free, Pro, Max (individual accounts)	Only if you allow it — but see the two exceptions below	Used "if you choose to allow us to use your chats...", if conversations are " flagged for safety review ", or on explicit opt-in.
Claude for Work (Team, Enterprise); and the API	No, unless you join a partner programme	"Anthropic does not use the data you share when using our commercial products to train our models , unless you choose to participate in our Development Partner Program."

Read first-hand on 13 September 2026 from: OpenAI Help Centre, *How your data is used to improve model performance* (updated 11 September 2026); OpenAI, *Enterprise privacy* (updated 8 January 2026); Anthropic Privacy Centre, *Is my data used for model training?* and *Does Anthropic Act as a Data Processor or Controller?* (both 17 March 2026). **These pages change. Check them, do not quote this table back at anyone in six months.**

The two holes in the opt-out

Suppose you have done the sensible thing and switched training off on your personal account. There are two ways your conversation can still be used, and neither is obvious.

Hole one — the thumbs-down button is a consent event

When an answer is poor and you click the thumbs-down icon, you are not filing a complaint into a void. You are submitting the conversation.

"Even if you have opted out of training, you can still choose to provide feedback... If you choose to provide feedback, *the entire conversation associated with that feedback may be used to train our models.*"

— OpenAI Help Centre, *How your data is used to improve model performance*

The other vendor is the same in substance, and adds a retention period:

"When you provide us feedback via our thumbs up/down button, we will store *the entire related conversation...* in our secured back-end **for up to 5 years."**

— Anthropic Privacy Centre, *Is my data used for model training?*

So the sequence that catches a careful person is this: you switch training off; you paste in a draft containing your client's name and the disputed sum; the answer is poor; you click thumbs-down to be helpful — and you have just submitted the whole exchange, with a five-year retention attached, by pressing the button that looked like quality control.

Hole two – flagged conversations, which you did not choose

One vendor states plainly that conversations flagged for safety review may be used and analysed **regardless of your training setting**:

*“Your conversations are flagged for safety review (in which case **we may use or analyze them** to improve our ability to detect and enforce our Usage Policy, **including training models** for use by our Safeguards team...)”*

— Anthropic Privacy Centre, *Is my data used for model training?*, listed as a circumstance in which chats are used

Think about which of your work is most likely to be flagged. A criminal brief. An allegation of fraud. A sexual-offence matter. A medical record in a compensation claim. The material a filter treats as sensitive is very often the material a lawyer handles — and it is precisely that material which is most likely to end up in front of a reviewer, or in a safety dataset.

THE PRACTICAL RULE, ON ONE LINE

On a personal account, treat the **thumbs-up / thumbs-down buttons as a submission**, not as feedback. If a conversation contains anything you would not put in an email to a stranger, do not rate it.

The part most lawyers get backwards

Here is the point that matters most, and it is the one almost nobody checks.

When lawyers think about a vendor holding client information, they reach instinctively for the familiar apparatus — *we are the controller, they are the processor, there will be a data processing agreement, the agreement will bind them*. Under Indian law the equivalent instinct is s.8(2) of the Digital Personal Data Protection Act, 2023: engage a processor “*only under a valid contract*”.

On a personal account, none of that apparatus is in place.

One vendor sets this out expressly. For its *commercial* products it says the customer is the “Controller” and it is the “Processor”, processing only on the customer’s instructions. Its consumer plans are dealt with in an entirely separate document, which says no such thing.

*“When a commercial customer creates a Claude for Work account (Team or Enterprise plan)... **the customer is the ‘Controller’** of the data submitted by its Users... **Anthropic acts as a ‘Processor’** of the data on behalf of the customer. Anthropic only processes the data as instructed by the customer in order to provide the Claude service.”*

— Anthropic Privacy Centre, *Does Anthropic Act as a Data Processor or Controller?*, 17 March 2026

The other vendor draws the same line from the other side: a data processing agreement is offered for **Business, Enterprise and the API**. Not for the personal plan.

WHAT THIS ACTUALLY MEANS FOR YOU

The personal subscription sitting open in your browser is not a processor arrangement. There is no data processing agreement. There is no set of instructions the vendor is bound to follow. You are a consumer using a service on its standard consumer terms — the same footing as someone asking it for a recipe.

The protections a lawyer assumes are in place are attached to a tier that lawyer is very often not on. **Moving from a personal plan to a business plan is not an IT upgrade. It is the step that creates the legal relationship you already believed you had.**

What the product is made of

Your paragraphs did not go to “the AI”. They went into a product with several separate parts, and it is worth knowing the parts, because each one is a different place where text can rest or be passed on.

One — the model

The weights. The next-word predictor. It reads what it is given, produces an answer, and retains nothing.

Two — the wrapper

Everything between you and the model: the screen, the hidden block of standing instructions the vendor adds to every request, the filters that inspect what goes in and what comes out, and the routing that decides which version of the model answers you today. **When someone says “we use ChatGPT”, this is the part they are describing.**


Three — memory

The model forgets everything between conversations. So every feature that appears to remember is built *around* it: your conversation history, re-sent with each message; the stored profile of things it has learned about you; the search index that looks through your uploaded documents and pastes relevant extracts into the request before the model sees it. **“It remembered our contract terms” is never a statement about the model.**

Four — tools

Web search, file upload, code execution, and connectors into a mailbox, a drive or a case-management system. **This is the part that reaches outward** — the point at which your information can leave the box you thought it was in.

FIGURE 1 – THE FOUR PARTS, AND THE TWO EVENTS



Training touched only the first part. What you type goes through all four. Parts two, three and four are often run with the help of other companies, in other countries — and both major vendors publish the list of who those companies are. Drawn generically; no vendor's internal design is depicted.

The list nobody reads

Those other companies are not a secret. European law obliges the vendors to publish them, and the lists are ordinary public web pages.

PROCESSOR · SUB-PROCESSOR — IN PLAIN TERMS

If you decide what happens to information, you are the one responsible for it. A company that handles it **on your instructions** is a **processor**. If that company hands it on to *another* company to help — a hosting firm, a support desk — that one is a **sub-processor**.

The point is simply that the chain does not stop at the company whose name is on your invoice.

OpenAI's list, last updated 9 July 2026, names nineteen outside companies and five of its own group entities. Hosting sits with Microsoft, Oracle, Google, Amazon, CoreWeave and Cerebras across more than two dozen named countries, India among them. Support sits with Intercom, Salesforce, Accenture and TaskUs. **Anthropic's list names twenty** — support with Intercom in the United States, Nutun in South Africa and Boldr in Canada; identity checks with Yoti in the United Kingdom.

Four things on those pages are worth two minutes of anyone's time.

What it is	What the page says	Why it matters to you
Switching on a feature	Merge API is “ <i>only used when the Customer chooses to enable Connectors in ChatGPT</i> ”	Ticking a box adds a company. A setting changed by any user quietly adds a firm processing across eleven countries.
Switching on a setting	Snowflake is in the chain “ <i>except where Zero Data Retention is used</i> ”	A contract setting removes one. Configuration is a legal fact, not an IT preference.
Reviewing flagged content	Flagged material may be sampled to TaskUs (Philippines), Accenture (US, Canada, Philippines), Cinder (US)	Real people, in other countries, may read flagged material — and sensitive matter is the most likely to be flagged.
One feature, two firms	Anthropic’s web search runs on Brave Search and TurboPuffer	A single toggle can add two companies , not one.

Both lists rendered first-hand on 13 September 2026: openai.com/policies/sub-processor-list/ · trust.anthropic.com/subprocessors

Where the law grips — and where it does not

Two regimes, two completely different approaches to that chain of companies. A European instinct applied in India will produce a confident, well-drafted, wrong answer.

Europe built a gate

*“The processor shall not engage another processor **without prior specific or general written authorisation of the controller**. In the case of general written authorisation, the processor shall inform the controller of any intended changes concerning the addition or replacement of other processors, thereby **giving the controller the opportunity to object to such changes**.”*

— Article 28(2), General Data Protection Regulation

Note what that is and what it is not. It is **not** a rule about who pays when something goes wrong. It is a rule about **who is allowed to join the chain, and your right to be told before they do**.

That rule is the only reason those public lists exist. A vendor cannot notify customers of changes without keeping a current published list and a channel to announce them on. Both vendors do exactly that, and both offer a subscription to the changes. Almost nobody has subscribed — so a right that is real and contractual simply expires unread.

India built no gate at all

This is a genuine absence, and it is worth stating precisely, because an absence stated precisely is more useful than a summary.

THE ENTIRE STATUTORY FOOTPRINT

The expression “**Data Processor**” appears **eight times** in the whole of the Digital Personal Data Protection Act, 2023 — once in the definition at s.2(k), once in an illustration, four times in s.8, once in s.6(6), once in s.11(1)(b).

Not one of those eight is a rule about a processor engaging someone else. No authorisation requirement, no duty to notify, no right to object. Neither the Act nor the Rules made under it says anything about the moment a fourth company joins the chain.

But no gate is not the same as no grip — and the provision that does the work is not in the obligations chapter at all. It is a right given to the individual:

“the identities of all other Data Fiduciaries and Data Processors with whom the personal data has been shared by such Data Fiduciary, along with a description of the personal data so shared”

— Section 11(1)(b), Digital Personal Data Protection Act, 2023

Read that against a list of nineteen or twenty companies. **A person in India can require you to name every company that has held their information.** Not a category — identities, with a description of what was shared. The only carve-out, in s.11(2), concerns sharing with another Data Fiduciary authorised by law in connection with offences or cyber incidents. Ordinary commercial sub-contracting is not excused.

And the Act does not stop at telling. On withdrawal of consent you must “*cease and cause its Data Processors to cease*” (s.6(6)). You must “*cause its Data Processor to erase*” what you gave them (s.8(7)(b)).

Cease. Erase. Name. Three instructions pointing down a chain, given to someone the Act hands no lever to pull.

FIGURE 2 — THE SAME CHAIN, TWO DIFFERENT LEGAL GRIPS

EUROPE · GDPR ART 28(2)

Approve before

- No new company **without your authorisation**
- You must be **told of intended changes**
- You get an **opportunity to object**
- **Result:** the public list exists

INDIA · DPDP S.11(1)(B) · S.6(6) · S.8(7)(B)

Disclose on demand

- **No** authorisation requirement in the Act
- **No** duty to notify, **no** right to object
- The individual may demand **every name**
- You must **make them stop, and erase**

Europe asks you to approve the chain before it forms. India asks you to disclose it on demand, and to reach down it and make things stop. Neither is possible for someone who has never read the list.

What to do on Monday

1. **Check which account you are actually logged into.** Personal or business. This single fact decides whether your text is used for training, and whether any processor relationship exists at all.
2. **Stop rating conversations on a personal account.** Thumbs-up and thumbs-down submit the whole exchange, opt-out or not, with retention of up to five years at one vendor.

3. **Open your vendor's sub-processor page and read it once.** Two minutes. You cannot answer a s.11(1)(b) request, or make anyone erase anything, without knowing who is on it.
4. **Check which features are switched on** — connectors, web search, file upload. Several add a named company to the chain. This is a setting anyone in the office can change.
5. **Subscribe to changes in the list.** The right to object runs from the date of notice. If nobody receives the notices, nobody exercises the right.
6. **If client material is going in, move to a business plan** — not for the features, but because that is where the data processing agreement and the processor relationship actually begin.

India Desk

Every duty in the Act that touches a processor is a duty to make someone else do something — and there are only three. Make them cease on withdrawal of consent (s.6(6)). Make them erase (s.8(7)(b)). Name them to the individual on request (s.11(1)(b)). All three point down a chain the Act gives the Data Fiduciary no statutory power to control. The drafting consequence is immediate: since the statute supplies no reach-down mechanism, the contract has to build one — a right to direct deletion, a right to a current sub-processor list, and a notification term — or all three duties are undischageable by design.

There is still nothing in the public record creating a transfer or sub-processing mechanism. Section 16 remains a power to restrict transfers to notified countries, and no country has been notified. Rule 15 of the DPDP Rules, 2025 reserves a power concerning personal data made available to a foreign State or its agencies, and no such order is in the public domain. Neither touches ordinary commercial sub-contracting. *Stated as an absence of record as at the date of this issue, not as a prediction.*

Running clock — computed 13 September 2026

The distances below are as at the date of this line and nowhere else in this issue. Everything above is in absolute dates, so it reads correctly whenever you open it.

13 November 2025 — DPDP Rules, 2025 notified. *304 days ago.*

20 February 2026 — IT Amendment Rules, 2026 in force. *205 days ago — in force.*

9 July 2026 — OpenAI sub-processor list last updated, as displayed. *66 days ago.*

13 November 2026 — Consent-Manager registration commences. *61 days.*

13 May 2027 — s.16 and Rule 15, and the substantive package including s.8 and s.11. *242 days.*

PRACTITIONER'S VERDICT

What to tell a colleague this week. The question is not whether the tool is safe. It is **which account you were logged into when you pasted the draft in**. On a personal plan your text may be used for training unless you turned it off; the thumbs-down button puts the whole conversation back in whether you turned it off or not; and there is no processor relationship and no data processing agreement of any kind — you are on consumer terms, the same as someone asking for a recipe. On a business plan those three things change together. **It is the cheapest compliance step available in this area and most of the profession has not taken it**, because it looks like a billing decision and is actually a legal one.

Where we are going

Next — Part 3: the memory layer. The model forgets everything between conversations, so every feature that appears to remember is built around it. That design is the reason deletion is contested rather than simple — and the reason “delete my chat” and “delete my data” are not the same request.

This issue is written for lawyers and is general in nature. It is not legal advice and no lawyer-client relationship arises from reading it.

Every statutory reference and every date in this issue was checked against at least two independent sources before it went out. Vendor terms were read first-hand on the dates stated. Where a claim could not be verified to that standard it does not appear, or it is stated as an absence of record rather than as a fact.

The archive, and everything else I have been building, is at wolfgangrush.github.io